

Practical Machine Learning

Steve Morin

November 19, 2016

Synopsis

This project will predict the manner in which exercise was done from a test data set. The model will be trained using labeled data and an evaluation will be done of the most accurate modelling method. The model that is generated will be used to predict the 'classe' in the test data set. Prediction outcomes will be evaluated by calculating the confusion matrix for the training/test data sets from the provided training data.

Load Libraries

```
## Include required libraries  
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(RColorBrewer)  
library(rattle)
```

```
## Rattle: A free graphical interface for data mining with R.  
## Version 4.1.0 Copyright (c) 2006-2015 Togaware Pty Ltd.  
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(randomForest)
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##  
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':  
##  
## margin
```

```
library(rpart)
library(rpart.plot)
options(warn=-1)
```

Load & Prepare Data

```
# Set the seed
```

```
set.seed(11111)
```

```
# Read testing data
```

```
testingData<-read.csv("D:\\Coursera\\Data Science Specialization\\Practical Machine Learning\\Assignment\\pml-testing.csv")
```

```
# Read training data
```

```
trainingData<-read.csv("D:\\Coursera\\Data Science Specialization\\Practical Machine Learning\\Assignment\\pml-training.csv")
```

```
# Partition the training data
```

```
trainIndex<-createDataPartition(trainingData$classe,p=0.6,list=FALSE)
```

```
# Subset the training data
```

```
trainDataPart<-trainingData[trainIndex, ]
```

```
# Subset the testing data
```

```
testDataPart<-trainingData[-trainIndex, ]
```

```
# Find columns with near zero variance
```

```
varNearZero <- nearZeroVar(trainDataPart)
```

```
# Remove columns identified as near zero from the training and testing data sets
```

```
trainDataPart<-trainDataPart[,-varNearZero]
testDataPart<-testDataPart[,-varNearZero]
```

```
# Find columns with greater than 75% NAs
```

```
columnsWithNA <- sapply(trainDataPart, function(x) mean(is.na(x))) > 0.75
```

```
# Remove columns with greater than 75% NAs
```

```
trainDataPart<-trainDataPart[,!columnsWithNA]
testDataPart<-testDataPart[,!columnsWithNA]
```

```
# Remove columns 1 - 5 from the train/test datasets as they are cannot be used in the analysis
```

```
trainDataPart<-trainDataPart[,-(1:5)]
testDataPart<-testDataPart[,-(1:5)]
```

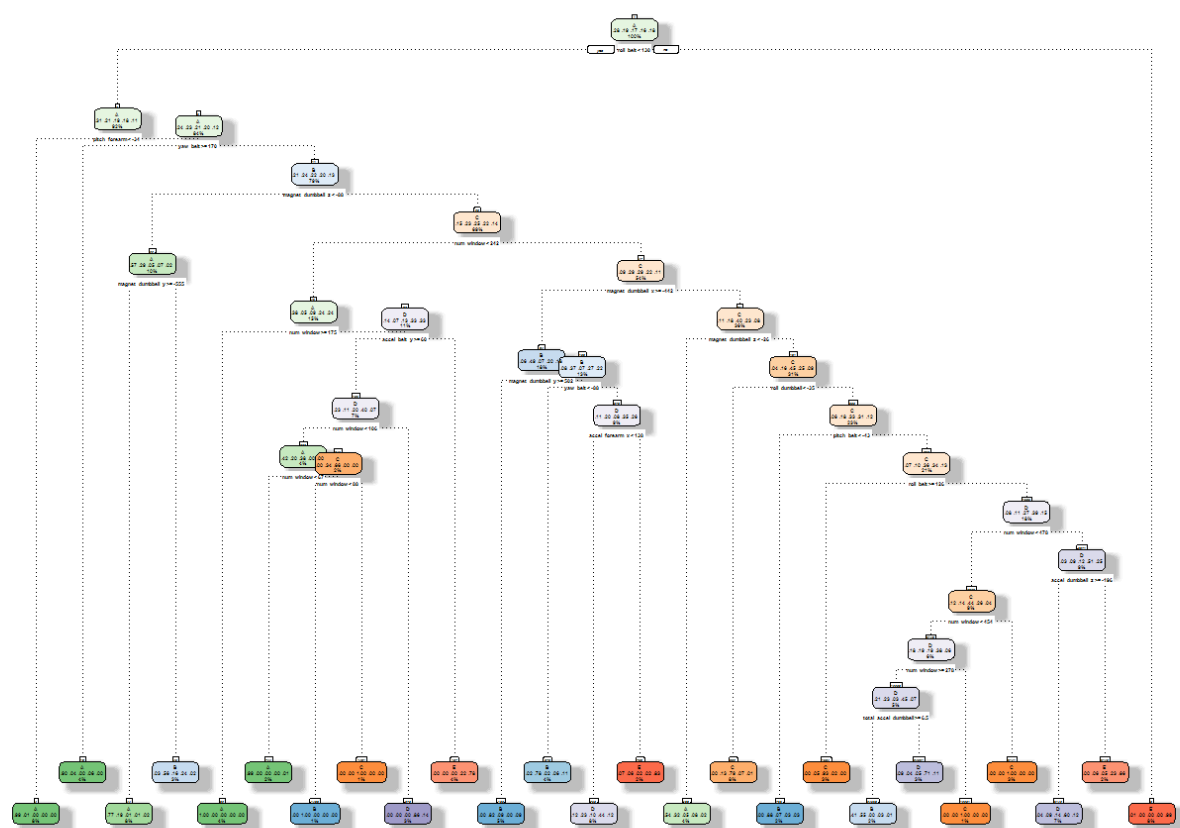
Decision Tree Method

```
# Create the Decision Tree
```

```
rpartModel <- rpart(classe ~ .,data=trainDataPart,method="class")
```

```
# Plot the Decision Tree
```

```
fancyRpartPlot(rpartModel)
```



Rattle 2016-Nov-19 16:16:13 morins

```
# Apply model to test data to obtain predicted classe
```

```
predictTestDataDT <- predict(rpartModel, newdata=testDataPart, type="class")
```

```
# Calculate the confusion matrix
```

```
confMatrixDT <- confusionMatrix(predictTestDataDT, testDataPart$classe)
confMatrixDT
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction   A    B    C    D    E
##           A 2006  264   21   67   25
##           B   82  940   85   97   90
##           C    1   82 1099   49    7
##           D  129  219  150  959  210
##           E   14   13   13  114 1110
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.7793
```

```
##           95% CI : (0.7699, 0.7884)
```

```
## No Information Rate : 0.2845
```

```
## P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##           Kappa : 0.7205
```

```
## McNemar's Test P-Value : < 2.2e-16
```

```
##
```

```
## Statistics by Class:
```

```
##
```

```
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.8987  0.6192  0.8034  0.7457  0.7698
## Specificity      0.9328  0.9441  0.9785  0.8921  0.9760
## Pos Pred Value   0.8418  0.7264  0.8877  0.5753  0.8782
## Neg Pred Value   0.9586  0.9118  0.9593  0.9471  0.9496
## Prevalence       0.2845  0.1935  0.1744  0.1639  0.1838
## Detection Rate   0.2557  0.1198  0.1401  0.1222  0.1415
## Detection Prevalence 0.3037  0.1649  0.1578  0.2125  0.1611
## Balanced Accuracy 0.9158  0.7816  0.8910  0.8189  0.8729
```

Random Forest Method

The second method tried was Random Forest using cross-validation for 3 iterations. Cross-validation is used to avoid overfitting the model and helps to make the model more generalizable to other data sets.

Random Forest Method

```
RFModel <- train(classe ~ ., data=trainDataPart, method="rf",trControl=trainControl(method="cv",
number=3, verboseIter=FALSE))
```

```
RFModel$finalModel
```

```
##
## Call:
##  randomForest(x = x, y = y, mtry = param$mtry)
##              Type of random forest: classification
##              Number of trees: 500
## No. of variables tried at each split: 27
##
##              OOB estimate of  error rate: 0.34%
## Confusion matrix:
##      A      B      C      D      E  class.error
## A 3346      1      0      0      1 0.0005973716
## B   10 2265      3      1      0 0.0061430452
## C      0      4 2048      2      0 0.0029211295
## D      0      0      8 1921      1 0.0046632124
## E      0      2      0      7 2156 0.0041570439
```

Prediction based on Test dataset with labels

```
predictTestDataRF <- predict(RFModel, newdata=testDataPart)
confMatrixRF <- confusionMatrix(predictTestDataRF, testDataPart$classe)
confMatrixRF
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 2232    9    0    0    0
##           B    0 1504    1    0    0
##           C    0    5 1367    8    0
##           D    0    0    0 1277    3
##           E    0    0    0    1 1439
##
## Overall Statistics
##
##           Accuracy : 0.9966
##           95% CI : (0.995, 0.9977)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9956
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity         1.0000   0.9908   0.9993   0.9930   0.9979
## Specificity         0.9984   0.9998   0.9980   0.9995   0.9998
## Pos Pred Value      0.9960   0.9993   0.9906   0.9977   0.9993
## Neg Pred Value      1.0000   0.9978   0.9998   0.9986   0.9995
## Prevalence          0.2845   0.1935   0.1744   0.1639   0.1838
## Detection Rate      0.2845   0.1917   0.1742   0.1628   0.1834
## Detection Prevalence 0.2856   0.1918   0.1759   0.1631   0.1835
## Balanced Accuracy    0.9992   0.9953   0.9986   0.9963   0.9989
```

```
# Perform final prediction based on Test dataset without Labels
```

```
predictFinalTestRF <- predict(RFModel, newdata=testingData)
predictFinalTestRF
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

Summary

The Random Forest method is more accurate than the Decision Tree method in this case achieving greater than 99% accuracy. As a result we used it to do the final prediction on the 20 original test cases.

End of Assignment

...